Which Factors Contribute to the Accuracy and Efficiency of Horse Racing Betting Markets?

Matthew Abel

King's College London

ABSTRACT

This is the abstract.

- A brief introduction to the objectives.
- A brief description of the main work.
- A brief description of the contributions, major findings, results achieved and conclusion.

Nomenclature

BSP	— Betfair Starting Price
CLOB	— Central Limit Order Book
HKJC	— Hong Kong Jockey Club
IP	— In-play (after the off)
PNL	— Profit and Loss
PP	— Pre-play (before the off)
WAP	— Weighted Average Price

1. Introduction

Horse racing markets have been academically studied as valuable research targets since at least the 1940s (Griffith, 1949). Racing markets are good to study for a few reasons. For one, they involve a large number of bettors. The bettors have varying levels of expertise ranging from absolutely zero to professional punters. These markets also closely mirror more ‘‘traditional’’ financial markets, and so research involving racing can be extrapolated out to the wider context of finance. Betting contracts are similar to options contracts in that they have a set expiry time and a payout determined by an external factor. The facets of horse racing make it an ideal ground to examine the efficient markets model (Snyder, 1978).

As of late, prediction markets have become a hot topic. Services such as Polymarket and Kalshi are offering the opportunity to bet on many world events. There has been research to

suggest that prediction markets are more accurate than other forecasting techniques including polls (Berg, 2008). Horse racing markets are also fundamentally prediction markets, and have a long history of widespread use. As a prediction market, horse races provide good data since there are so many and the approximate date and time is known before hand. It is easier to determine their accuracy in this context as there are many repeated similar markets, unlike predictions for political events which can be very far apart.

Existing research into

1.1. Motivation

I was motivated to research this area due to the recent popularity of prediction markets as a tool for forecasting. As well as this, the regulation surrounding prediction markets is currently heavy within the UK. This is because they fall under the jurisdiction of the Gambling Commission. Many of the newly popular prediction market sites are currently illegal to use within the UK, but established exchanges such as Betfair are completely legal. Another interesting factor of being classed as gambling means that winnings within the UK are not taxed. This means that an equivalent strategy within a gambling market and the stock market may prove ultimately more profitable in the gambling market as the only fees to be paid are commission to the exchange. In the case of Betfair, commission can be even less than just 2% of successful bets.

1.2. Scope

This project observes a variety of different classifications of horse races. Each classification type is examined to see if there is any link between the classification of a horse race and it's related betting market's accuracy and efficiency. In order to achieve this, a database of horse races is compiled using data from the Betfair exchange.

1.3. Aims

The primary aim of this project is to determine which factors are most significant in influencing the accuracy and efficiency of horse racing markets. There are many different ways to classify horse races including but not limited to time, distance and race type (Hurdles, Chase, Handicap, etc.). Additionally, races can be classified by their market factors such as trade volume or overround.

Accuracy of the market is a measure of how well the listed odds represent the true odds of a horse winning. In a market that was 100% accurate, the odds given would provide a 0% return on all winning bets. To determine the accuracy of the odds, I will be comparing various test statistics such as the Brier score and Log-Loss.

An efficient market is one in which prices are said to ''fully reflect'' all of the available information and are shown to be *weak-form* efficient when this can be demonstrated using an information set made exclusively of historical price data (Fama, 1970). A systematic betting strategy using historical pricing data that is able to generate a consistent profit demonstrates a lack of *weak-form* efficiency within a market (Hausch, 1981). In order to determine which groups of races are *weak-form* efficient, I shall apply various betting strategy profiles to them in isolation. Some of these strategies will simple and some will be adapted from previous work.

1.4. Objectives

- Perform a systematic review of existing research to determine existing standards for market accuracy and efficiency. Identify whether racing markets are generally considered to be accurate or efficient.
- Obtain an appropriate dataset for research.
- Categorise the data into potentially useful groupings.
- Examine each grouping to determine whether or not they fit the previously identified definitions of accurate or efficient.

2. Background

2.1. Bookmaking

2.1.1. Fixed-Price Odds

A fixed-price bookmaker will quote prices (odds) at which they are willing to sell bets at. If the bet pays out, the bettor receives the return originally quoted to them. The bettor is given a guarantee of their winnings before placing the bet. It is important to note that it is unlikely a bookmaker will allow a bettor to ‘‘lay’’ a bet, that is to bet that something *will not* happen. Bookmakers make their money by appropriately setting odds such that in all cases (or at least over a period of time), they will ensure a profit. This is done by changing odds to account for the amount of money bet already (balancing the books) and by providing unfair odds that do not sum to one.

2.1.2. Parimutuel Odds

A tote (otherwise known as parimutuel betting) is a system in which money for all bets is placed into a pool and then distributed to winning bettors. This means that bettors are not given an exact payout quote at the time of purchasing their bet, but they may be given an estimate. Tote operators make money by taking a certain percentage of the pool for themselves, typically called the ‘‘track take’’. These types of bets are typically facilitated by the track operator themselves, such as the HKJC (Hong Kong Jockey Club, 2026).

2.1.3. Betfair Exchange

Betfair is a UK based company which operates a betting exchange service. A betting exchange is different from a typical bookmaker in that instead of the operator quoting odds prices or offering a tote service, it allows users to form bets with one another. Betfair does this using a CLOB, allowing users to either back (long) a bet or lay (short) one. Within the context of a horse race, this means that the win market (where bets are payed out only if the horse places first) features a different asset for every selection (runner) in the race. They can be thought of as largely analogous to options contracts, where the layer agrees to pay the backer a certain sum of money dependant on the outcome of the race.

A unique feature of the Betfair racing market is its "starting price" betting. Bettors who use the starting price option have their bets matched into the market at the off (when the race begins). They are also able to set minimum odds to back or maximum odds to lay. The price determination maximises volume and gives backers and layers the same price. In order to do that, existing

positions within the CLOB may be matched against as well. Starting price bets aren't guaranteed to be matched if the odds limit is at or above the starting price (Betfair,).

2.2. Literature Review

2.2.1. Accuracy of Prediction Markets

2.2.1.1. Favourite-Longshot Bias

The favourite longshot-bias has been demonstrated within racing markets as far back as 1949 (Griffith, 1949). This is a phenomenom in which bettors tend to prefer to bet on outlier horses, as opposed to the safer options. In turn, this moves the odds given by the market as favourite horses are given worse implied odds and longshots are given better implied odds. Other than the long and short-odded horses, Griffith found that the odds for horses were on average accurate.

Some studies have attempted to determine why this is the case. Weitzman (1965) used utility theory to derive a utility curve for one bettor and found it to be similar to the theoretical one proposed by Markowitz. To follow this study, Ali (1977) observed over 20,000 races and found the favourite-longshot bias to be present. Ali creates a wealth utility function and determines that the average bettor values risk, thereby causing the discrepancy in prices. Subsequent research has also conjectured that bettors have a preference for risk at the racetrack. Kanto et al. (1992) found this to be the case using a data sample from Finland.

Golec and Tamarkin (1998) suggested an alternative explanation for the favourite-longshot bias. They claimed that it was possible that bettors could actually be risk-averse and instead have a preference for skewness. To back this up they examined popular betting strategies and determined that they were better aligned with a preference for skewness rather than risk. Another theory is put forward by Snowberg and Wolfers (2010) which states that the bias can be accounted for by the misperceptions of risk by the bettors.

The bias has also been shown to not be universal. In contrast to the evidence found in North American markets, Busche and Hall (1988) found that bettors at the Hong Kong Jockey Club do not exhibit the same apparent risk-loving behaviour. They posit that the larger betting volume in Hong Kong may influence the bias.

2.2.1.2. Use of Markets for Forecasting

2.2.2. Efficiency of Prediction Markets

- efficient market theory
-

2.2.3. Strategies for Horse Race Betting

3. Main Results

3.1. Objectives, Specification & Design

3.2. Methodology & Implementation

3.2.1. Data Acquisition

3.2.1.1. Betfair Exchange UK Horse Racing Win Markets

Historical odds pricing data is typically very difficult to find if not impossible from most bookmakers. This was a signifcant factor in choosing to use prices from the Betfair exchange. Betfair provides historical betting data at multiple levels of resolution including a generous free tier. The data provided in this set is structured like exchange data, with multiple entries per race documenting changes in prices over time. For my purposes, this is beyond what I need as I am only interest in the BSP and not the price evolution. I used this dataset at first, but found one better suited to my uses.

Betfair also provides a directory of starting price files and allows users to filter them by location and market type. This dataset contains some useful information not included in the free tier of the standard historical data files such as trade volume and minimum and maximum prices for each selection. This dataset is also significantly larger, containing data from as far back as 2008, compared to 2016 with the other set.

I produced a Python script to automatically download all of the UK win market files and process them to be used within Pandas. The script saves the data to binary files for quick retrieval. It makes sure to not redownload files it has already obtained so it can easily update the existing database. As well as this, it automatically slows down when rate limited to ensure that all files are downloaded.

The final dataset contains over 170,000 races from 24,000 different events in the UK, starting from September 1st 2008. There are over 1.5 million entries for selections within the races. This data is formatted into three different tables: events, races and, runners.

3.2.1.1.1. Events

Field	Data Type	Notes
`event_id`	`Integer`	Unique identifier provided by Betfair
`location`	`String`	
`event_date`	`Datetime`	

3.2.1.1.2. Races

Field	Data Type	Notes
`race_id`	`Integer`	Unique identifier provided by Betfair
`race_type`	`String`	Indicates the format of the race: handicap, hurdles, chase, etc.
`race_time`	`Datetime`	
`distance_f`	`Integer`	Furlongs are one eigth of a mile long.
`race_no`	`Integer`	Races are given numbers indicating their sequence position within an event.
`event_id`	`Integer`	Relational key.

3.2.1.1.3. Runners

Field	Data Type	Notes
selection_id	Integer	This identifier is unique per horse, not per entry.
selection_name	String	
win_lose	Boolean	Win=True, Lose=False
bsp	Float64	Betfair Starting Price
implied_odds	Float64	1/BSP
ppwap	Float32	Pre-play weighted average price.
morningwap	Float32	Weighted average price of bets before 11AM GMT
ppmax	Float32	Pre-play max price.
ppmin	Float32	Pre-play min price.
ipmax	Float32	In-play max price.
ipmin	Float32	In-play min price.
morningtradedvol	Float32	Trade volume before 11AM
pptradedvol	Float32	Trade volume before the off
iptradedvol	Float32	Trade volume after the off
race_id	Integer	Relational key.

3.2.1.2. Hong Kong Jockey Club

I also downloaded data from the website of the Hong Kong Jockey club in order to compare it against the UK market. This was done by scraping their webpages for historical racing events. Ultimately it proved that this data was not particularly useful, however. The available data I was able to collect only went back to 2022 and the reduced number of racecourses operated by the HKJC meant that the number of races was only around 2000, compared to the 170,000+ of the UK dataset.

3.2.2. Software Development

I decided to use python through Jupyter Lab to write and run my software as it is what I had the most experience with for statistical analysis.

3.2.2.1. Data Structure and Statistics

3.2.2.1.1. Pre-Processing

I pre-processed my data to include some more fields that would be useful. Pre-processing fields like this means that the calculations do not need to be repeated later on.

The first field to be added was a normalised probability field. This is useful as the implied odds given by the starting price (almost) never sum to be one. The sum of implied odds given is known as the ‘‘overround’’. An overround greater than one for a back market shows that there is negative arbitrage and a weighted selection of back bets is guaranteed to lose money. Since the data for this project is using the BSP, it can trivially be shown that it is possible to arbitrage almost every single market. This is beacuse for BSP, both back bettors and lay bettors are given the same price. In real terms this arbitrage is actually impossible to achieve however, as the BSP is not possible to know before the off and by betting this way you are certain to move the market.

I then added fields for partial pre calculation of the Brier score and Log-Loss. This is particularly useful as it means that to determine the scores for a certain subset of races, they can simply be summated using the groups generated in the next section.

3.2.2.1.2. Categorisation

Given the available data from my dataset I decided upon a list of ways to categorise the horse races: race number, distance, hour of day, pre-play trade volume, race type (handicap, hurdles, etc.), location, number of runners and, overround. To do this I produced a list containing these different classifications of races created using the `groupby` method in Pandas. The `groupby` method makes it easy to perform aggregate calculations per group created through the method. This means that a function such as `sum` will be ran per group, instead of for all data in the DataFrame.

3.2.2.2. Trading Strategies

Trading strategies are implemented using a system in which the strategy function recieves a list of available runners to bet on and returns a ‘‘bet slip’’. The bet slip contains information about which race and runner it chose, as well as the amount bet and whether or not the selection won. The amount bet is able to be submitted as either an absolute value, or as a percentage of the currently held bank value.

The bet slips are passed to a function which actions the bets and determines the PNL for each one. It also sorts the bets over time to generate a timeline of events.

Each betting strategy is written using vector based functions. This allows the underlying library (Pandas or Numpy) to perform many operations in parallel. There are two primary benefits to this. Firstly, it significantly improves the readability of the code as the need for looping is removed. The code is much more representative of what is actually happening. Secondly, there is a major performance increase by writing the software in this way.

3.2.2.2.1. Naive Strategies

The first naive strategy simply bets one unit of currency on the favourite (runner with the highest implied probability) in every race. Due to how I set up the strategy implementation, this is incredibly simple as it just involves grouping the runners per race and then submitting a bet of one. If the favourite horse is given odds lower than its real odds, then this strategy should consistently make money. Breaking even is an implication of the favourite being accurately priced over time.

The other naive strategy is similar to the first, but instead of betting a constant amount on each favourite, it attempts to ‘‘hedge’’ the bet by betting lower for worse odds and higher for better odds. If the prices for the favourite horse are accurate, then this strategy should consistently break even.

3.2.2.2.2. Learning Strategy

I created a very simple learning strategy that uses a rolling average to estimate the current likelihood of the favourite horse winning. It assumes the current average of win rate of the favourite horse to be the true win rate and uses the Kelly Criterion to size a bet. That is: $f^* = (bp - q)/b$ where f^* is the fraction of wealth to bet, p is the likelihood of the bet winning, q is $1 - p$ and b is the proportion of wealth won if successful (Thorp, 2006). It is worth noting

that this does not account for the impact that a bet will have on the market. This is a very simple strategy which does not use any information outside of the historical prices. If this strategy is able to make a profit, it indicates that the market is not *weak-form* efficient.

3.3. Results, Analysis & Evaluation

3.3.1. Accuracy

3.3.2. Efficiency

4. Legal, Social, Ethical & Professional Issues

- Use of live animals?
- Profit Derived from gambling losers
- Similarity to stock market
- legality of companies betting

5. Conclusion

References

Ali, Mukhtar M., “Probability and utility estimates for racetrack bettors,” *Journal of Political Economy* **85**(4), pp. 803-815, University of Chicago Press (1977).

Berg, Joyce E.,Nelson, Forrest D., Rietz, Thomas A., “Prediction market accuracy in the long run,” *International Journal of Forecasting*, US Presidential Election Forecasting **24**(2), pp. 285-300 (2008-04-01).

Betfair, *Betfair SP FAQ*, Betfair. Available at: https://promo.betfair.com/betfairsp/FAQs_detailed-Workings.html [Accessed Mon Jul 13 2026].

Busche, Kelly & Hall, Christopher D., “An exception to the risk preference anomaly,” *The Journal of Business* **61**(3), pp. 337-346, University of Chicago Press (1988).

Fama, Eugene F., “Efficient capital markets: a review of theory and empirical work,” *The Journal of Finance* **25**(2), pp. 383-417, [American Finance Association, Wiley] (1970).

Golec, Joseph & Tamarkin, Maurry, “Bettors love skewness, not risk, at the horse track,” *Journal of Political Economy* **106**(1), pp. 205-225, The University of Chicago Press (1998).

Griffith, R. M., “Odds adjustments by American horse-race bettors,” *American Journal of Psychology* **62**(2), pp. 290-294, University of Illinois Press (1949).

Hausch, Donald B.,Ziemba, William T., Rubinstein, Mark, “Efficiency of the market for racetrack betting,” *Management Science (pre-1986)* **27**(12), p. 1435, Institute for Operations Research and the Management Sciences, Linthicum, United States (1981-12).

Hong Kong Jockey Club, *Pari-Mutuel Local Pools*, Hong Kong Jockey Club (2026). Available at: https://special.hkjc.com/e-win/en-US/betting-info/racing/beginners-guide/local-pools/ [Accessed Wed Jul 08 2026].

Kanto, Antti J.,Rosenqvist, Gunnar, Suvas, Arto, “On utility function estimation of racetrack bettors,” *Journal of Economic Psychology* **13**(3), pp. 491-498 (1992-09-01).

Snowberg, Erik & Wolfers, Justin, “Explaining the favoritelong shot bias: is it risk-love or misperceptions?” *Journal of Political Economy* **118**(4), pp. 723-746, The University of Chicago Press (2010).

Snyder, Wayne W., “Horse racing: testing the efficient markets model,” *The Journal of Finance* **33**(4), pp. 1109-1118, [American Finance Association, Wiley] (1978).

Thorp, Edward O., “The kelly criterion in blackjack sports betting, and the stock market” in *Handbook of asset and liability management* **1**, pp. 385-428, Elsevier (2006).

Weitzman, Martin, “Utility analysis and group behavior: an empirical study,” *Journal of Political Economy* **73**(1), pp. 18-26, University of Chicago Press (1965).

Table of Contents

Nomenclature . . . 1
Introduction . . . 1
  Motivation . . . 2
  Scope . . . 2
  Aims . . . 2
  Objectives . . . 3
Background . . . 3
  Bookmaking . . . 3
    Fixed-Price Odds . . . 3
    Parimutuel Odds . . . 3
    Betfair Exchange . . . 3
  Literature Review . . . 4
    Accuracy of Prediction Markets . . . 4
      Favourite-Longshot Bias . . . 4
      Use of Markets for Forecasting . . . 4
    Efficiency of Prediction Markets . . . 4
    Strategies for Horse Race Betting . . . 4
Main Results . . . 4
  Objectives, Specification & Design . . . 4
  Methodology & Implementation . . . 5
    Data Acquisition . . . 5
      Betfair Exchange UK Horse Racing Win Markets . . . 5
        Events . . . 5
        Races . . . 5
        Runners . . . 6
      Hong Kong Jockey Club . . . 6
    Software Development . . . 6
      Data Structure and Statistics . . . 6
        Pre-Processing . . . 6
        Categorisation . . . 7
      Trading Strategies . . . 7
        Naive Strategies . . . 7
        Learning Strategy . . . 7
  Results, Analysis & Evaluation . . . 8
    Accuracy . . . 8
    Efficiency . . . 8
Legal, Social, Ethical & Professional Issues . . . 8
Conclusion . . . 8
References . . . 9